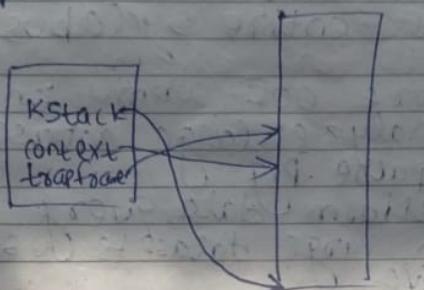


Prof. Rebero - IIT Chennai? Both use
IIT Delhi JVVG

All exams will be online.

* struct p80c



ushort → 2 bytes

* struct trapframe :

→ The lowermost esp and ss
all the pointers to older application
stack

* All read write codes are present
in vsys.S.

After INT instruction hardware does
some pushes & jumps control to
vector64. The stack change happens
OSL → kernel process stack happens
in hardware.

* vint esp in trapframe is the push \$0
in vsys.S

* See argstr() algint() in sys_codes
4+4xh → calling convention

Threads & Signal

22/2/2022

Thread - independent flow of code execution

→ fact() divided into 2 parts:
 $1 \rightarrow n/2, n/2 + 1 \rightarrow n$

thread-join is like wait. It waits for each thread to be completed.

* gcc - optimised

* Thread - flow of code which uses CPU

* Concurrency: progress at same time
Parallel: execution at same time

For one processor, parallelism not possible.

Parallelism needs multiprocessor

Scheduler also schedules threads.

* fork() → heavy-weight. As PCB needs to be setup.
thread creation takes less effort

* YouTube keeps running on 1 tab due to multi-threading.

→ when we run a function on a thread that function may call another fn. so each thread needs a separate stack.

* we need to do a context switch (Save registers) to switch from 1 thread to another.

* code of threads is code of process

* Responsiveness - Even if 1 thread blocks (scanf()), then other threads can keep executing.

CPU * we don't need separate shared mem while using threads.

for * For multicore systems each core has different interrupt (eg timer)

not * xv6 gives us only processes.

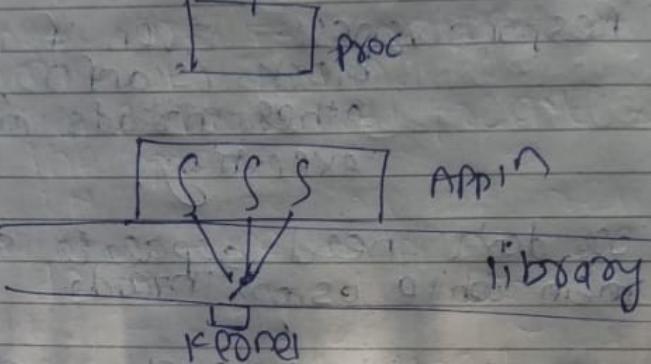
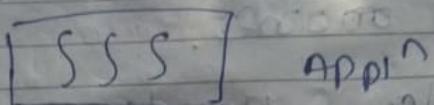
ng * user threads - The applications handle them

s. * How to do scheduling in userland?

for * If OS doesn't provide threads, it is possible that userland lib can create & manage threads.

ed * Even if kernel gives us threads, a user library has to map ~~these~~ user threads → kernel, otherwise app codes cannot access those threads.

- * If there are no kernel threads



→ If we have no kernel threads
then kernel thread means process

- * Pthread on linux is 1:1 mapped

* In Many → one, if 1 thread
does blocking entire process
is blocked. ~~as kernel~~
understands only processes

* Disadvantage of one-one thread:
For one user thread one kernel
thread is needed. Hence more
resources.

- * Two level model - Hybrid.

* `(clone())` → creates a kernel
thread
Linux → understands only threads!

By
JSP

flags argument to clone() gives choice of what part of process needs to be duplicated.

fork() is a fn on top of clone().

- * Learn clone() to understand threads.
- * Follow syntax of pthread lib for doing project as it is POSIX standard.
- * what happens when fork() or exec() are called in a thread
- * Thread cancellation - terminate a thread

main() {
 f();
}
f() {
 pthread_create(&f, NULL, &clone, f);
 clone(f);
}

clone() {
 Linux Kernel
 By default
 duplicated stack

* If for every pthread_create()
we call clone(), then it's
1-1 mapping.

* How to do many-one mapping

This one is created at the
time of fork().

So one flaw is present, which
is the process.

For many-one mapping we have
to do job without calling clone.

* We need struct thread &
char stack[].
context :-
fn pointer ; -

→ context → depends on calling convention.

pthread_create(fn)

t = malloc (thread)
threads[i] = t
switch ()

In linux we have userland lib
fns which get context & switch
the context

→ First write project using existing
lib fns.

* Change context → change stack.

Q → what is the context

switch() → take context of 1 thread, put on context(stack) of library, then switch to new stack

* See fns getContext(), setContext(), makeContext(), longjmp(), setjmp().

* To get more marks - write
getcontext(), setcontext()

24/1/2022

Signals

Page No.

Date / /

Segment - event

synchronous - predictable

System calls to make signals -

synchronous way

signal() → system call

signal.h

SIG_INT - Macro for signal

Kernel has default signal handling

SIGINT occurs when we press

ctrl + C

ctrl + z - SIGSTOP

* Synchronous way of sending
Signal - kill()

→ Kill is not always used to
kill a process.

So one process can send signal
to another. This is synchronous
way

* man + signal

* SIGKILL is 9.

Kill -9 pid will DEFINITELY
kill the process

Kill -9 SIGINT pid

→ sends signal to pid.

shows all syscalls that command is making

* strace kill -s SIGINT 8179

* kill -s SIGSTOP pid

fg → brings stopped process into foreground
i.e. process was paused → it has continued.

SIGCONT = continue

* signal() is no longer preferred choice.

sigaction() is POSIX standard

SIGCHILD → send by kernel to parent when child dies

Imp] SIGALRM → Timer!

* Signal handlers can only run when process is scheduled
The scheduler must now check if any pending signals are there
How to jump to signal handler from scheduler?

* Signal handlers doesn't work acc. to LIFO stack

* Thread pool
→ normally when we create a thread it starts running immediately with a fixed context

→ Thread pool is created so that it is ensured that threads will be available when they will be needed.

* global - visible in all files
static global - visible in only this file.

Hence some fns are also marked as static.

* --Thread is specific to gcr. Not

* one-one mapping : Huge resources required

* How does app? know that kernel threads are blocked?

→ Kernel has to communicate with user

Ucalls → backward call, kernel calls user code.

* How does a kernel call user fn?

* LWP → data structure given by kernel to user library

LWP part of Kernel DS

User can access LWP

so user library is a wrapper around LWP

* Linux doesn't distinguish b/w process & thread

* start / fork program

→ see that it internally calls clone

this

marked

Not c!

once

kernel

the

el

fn?

on by
very

already

process &

1 Mar 2022

O.S. XV6 - paging code

We started with bootasm.s →
bootmain.c → entry.s

entry creates a page table with a
single entry.

Go to main.c:

(1) kinit1 (and P2V ($4 * 1024 * 1024$))

Physical to
Virtual
↑
4 MB

adds KERNBASE

→ We can see 'end' in kernel.sym.
'end' is last address of ELF file.

Range between vstart & vend is less
than 4 MB.

Ignore the lock.
Go to free_range()

PGSIZE is 4K

PGROUNDUP - Rounds up addr to
a page boundary (4K,
16K, etc...)

so for every address on a page
boundary, kfree() is called.

kfree() is adding a chunk to
a linked list.

kfree() is passed a ptr which is
aligned on a page boundary

Up to now kernel code uses only
22.4 MB - PHYSSTOP

panic() is a fn which will make the
kernel stop & enters in an infinite
loop.

- * Kmem is global structure
- * A certain part of the node of
linked list (4 byte ptr) is used to
link
- * Only the free chunks are chained
together
- * Rn: logical address = Physical address
bcz segmentation setup is done.
Any address $2GB +$ is mapped to
 $0 \rightarrow 4MB$
- * Each page in xv6 is 4KB
- * So kinit() creates a freelist.
Why free pages created? \rightarrow Bcoz all
next fns need this stop
- * kinit2(P2V(4MB), P2V(PHYSSTOP))
- * Rn we are living in a world where
there are 3 entries in page
directory. $0 \rightarrow 4MB$ only
Hence kinit1(end, 4MB)
max is 4MB only

- * `kvmalloc()` → does mem setup for scheduler
- * `lcr3` → present in `x86.h`
In `x86.h` → there are wrappers around assembly code
 $\text{CR3} \rightarrow$ base of page directory
- * `setupkvm()` returns new kernel page directory address.
- * The `lcr3()` won't get called.
It is basically an inline fn.
Its code is copy pasted directly
- * `setupkvm()` is called from many places.
- * In `kvmalloc` we are going to map certain regions of logical memory to certain regions of physical memory
- * In xv6:
 - For process: $0 \rightarrow 2\text{GB}$
 - For Kernel: 2GB onwards
- * `PHYSSTOP` → limit on physical memory currently (224 MB)
- * `kmap[]` tells that `KERNBASE` should be mapped from 0 to `EXTMEM`

- * ending virtual address
= physical end = physical start + virtual
- * come back to setup VM

Go to Kalloc() → It gives us a 4K frame.

Kalloc() picks up first frame from freemem list

kfree() → puts a frame on free frame list

- * Remember PDE and PTE from x86 mem management.

Every PDE entry points to a page table

- * memset(pgdir, 0, 4K)
as everything is set to 0
the PS bit in PDE is also 0.
So it says that it's 4K page

- * mappages() → creates entries within page dir & page table

mappages() job is to set up entries in page directory & also the page tables. 2 way paging

~~mappages~~ Go to mappages

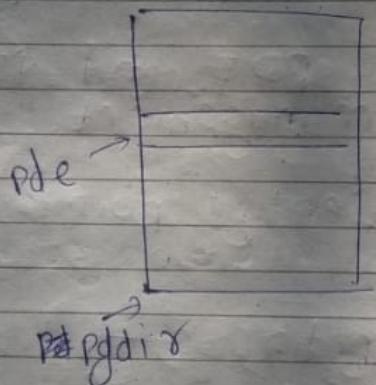
PGRONDOWND(a) aligns address downwards on a page size

virtual addr is a
physical addr is pa

- * walkpgdir() returns a pointer to a page table

If alloc = 1 then creates new page tables

PDX() → gives Page Directory Index (10 MSB bits)



- * pde → particular entry

* Entries in a Page directory & page tables are PHYSICAL address

PTE_ADDR → gives address of page table entry

- * Initially *pde & PTE_P is false as we haven't set it to 0

Go in else part

pgtab takes 1 page using Kalloc()

Now we will set up Page Table

- * walkpgdir() checked if va existed in pgdir.
- * After kvmalloc() in main we switched from 4 MB pages to 4K pages & ~~then~~ switched to 2 level paging.
- * After kvmalloc() we can access memory upto PHYSSTOP.
kinit2() uses 2nd arg as PHYSSTOP
- * Goto seginit() from main
see struct CPU

In seginit() again mapping is same 0 → 4GB

only diff. is that here UCODE has privilege as DPL_USER
UDATA has privilege level DPL_USER

so now the processes will use only the UCODE & UDATA

2/3/2022

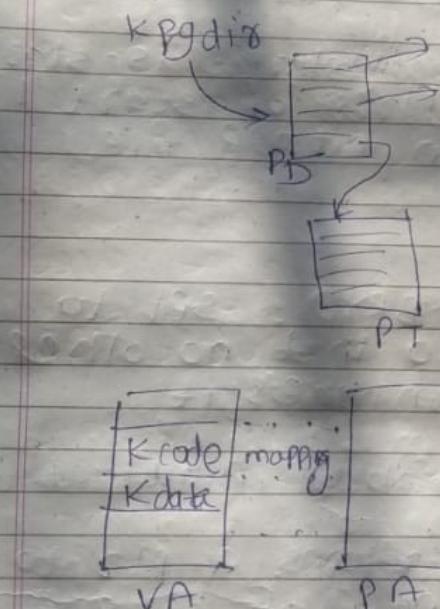
O.S.

kinit1
kinit2 } create free frame list

Kalloc, Kfree & managing free frame list

KVmalloc → setupKvm

setupKvm():



Q. Can the userinit() be moved to line 35?
AS userinit() starts first user process

* Goto userinit() → creates init process by hand

extern char _binary_initcode_start

→ This _binary_init_code can be seen
in kernel.sym

→ initcode.s, see in Makefile.

* We are going to create a
process called introde.
The initcode creates init.

* _binary_initcode_start : is starting
address of start fn
in initcode file.

* allocproc():

p → state = EMBRYO : set to
embryo so that no other
process can access it

nextpid is global initialized to 1

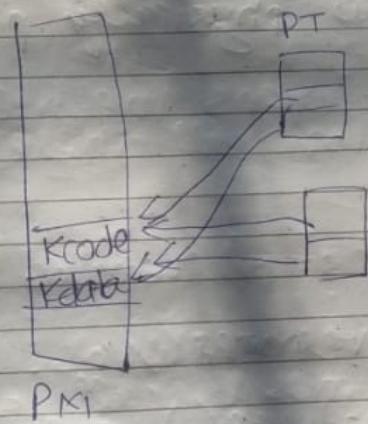
Kalloc() → returns a ptr to a
page of size 4 KB

* We have to create a trapframe
rn bcoz once we return from
kernel we want to simulate
returning from a trap

* trapset fn is present in trapasm

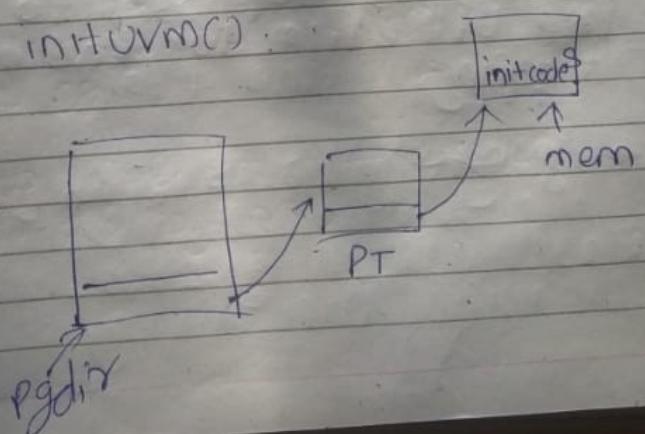
* Draw diagram of what allocproc()
is doing

- * init proc is global pointer to init process
- * setupKvm() maps kernel code & kernel data & creates page tables
 - The kernel code & data already exists in Physical mem. only the mapping is done in setupKvm



- * initUvm → init uses virtual memory
- * mapPages() → create page table mappings.

* IN initUvm():

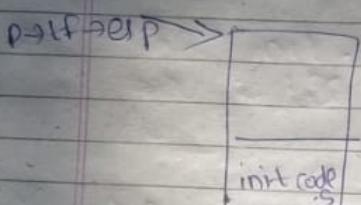


$p \rightarrow sz$ is PGSIZE, But actually size stored is 0x002C.

* SEGCODE & SEG_UPDATA all 3 and 4 broz in gdt the 3rd & 4th entries correspond to user code & data

* $p \rightarrow tf \rightarrow esp = PGSIZE$

so esp now physically points to

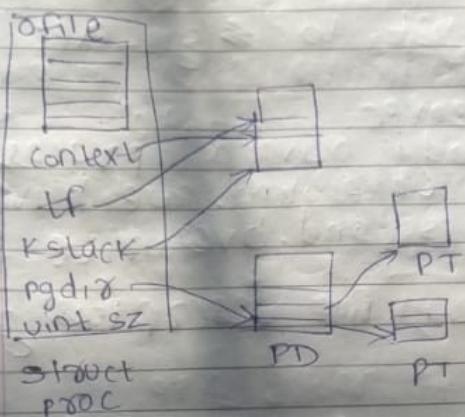


* SKIP $p \rightarrow cod = namei("/", "/")$

* After scheduler() schedules switch() on initcode all goto forkret

03/ Mar/2022 O.S.

Code of exec :

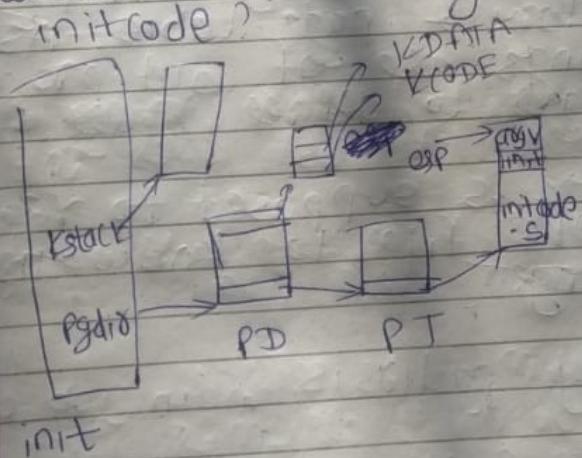


* Goto initcode.S

→ push \$argv
push \$init
push \$0

→ See what is argv and init

Where is this pushing happening in initcode?



→ From initcode.S we made int 64
So now stack changed.

Now ESP points to KStack

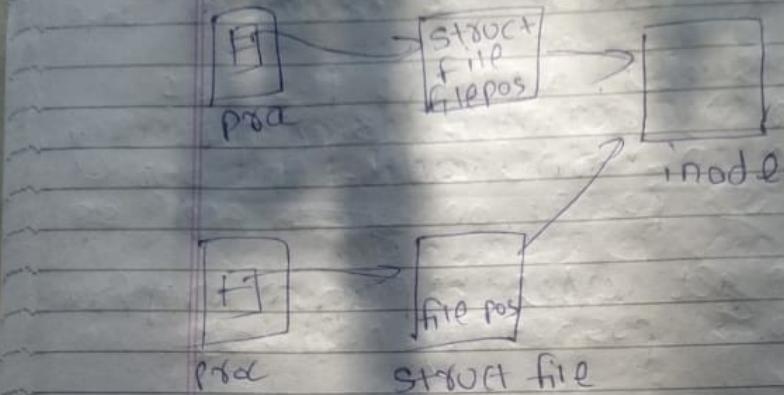
How do we get to sys-exec?

→ Goto sysexec.

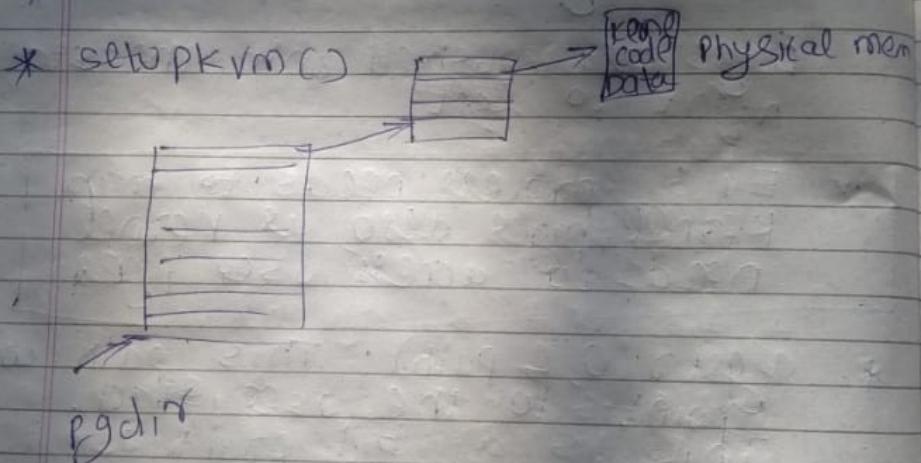
→ We already have a process which is executing. We have to clear this and overwrite everything.

- * On a fork, file descriptors get copied.
 - * The init.c opens the files 0 1 2. So after fork all these will be duplicated.
 - * So, on exec: pid remains same.
 - * exec() never returns. So exec() must call sched().
 - * We have to change the page directory & page table. As the mappings are of old process. We have to kfree() all frames.
- exec() is supposed to read the ELF file to know the mappings.
- * Changing pgdir is main job of exec.
 - * Read code of sys-exec.

- * see how exec() is called in sh.c.
- * In syscalls we have fns like argstr(), argint(). Bcoz currently we are in kernel stack. And our args are pushed on ~~system~~ appn stack
we can get older esp from trapframe
This was pushed by hardware
- argstr(0, &path) means fetch argument in path
- Args are pushed onto the stack right to left
- * Each process needs to map kernel pages also as kernel needs to access user pages
- * argstr(0, &path): This path will point to the user stack.
This is a bad practice.
- * exec() does the actual work.
- * * ignore begin-op()
- * namei() finds file on the disk & creates an inode data structure
inode is unique rep of file.



- * `readi()` → reads file



- * `map_pages()` doesn't allocate memory! It just does the mapping

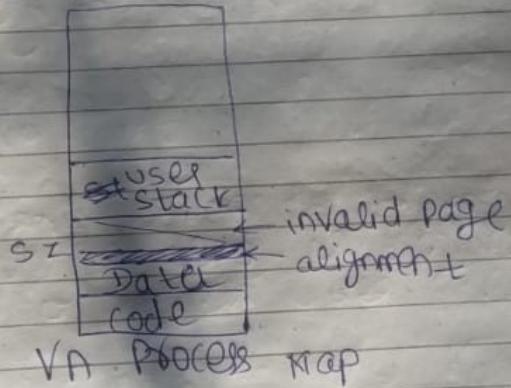
- * The offset in elf file is byte-wise offset in the file

- * `allocuvvm()` is called
→ `allocuvvm()` assumes that process is of some size & increases the memory

loaduvm() → loads file contents into memory allocated by allocuvm()
ip → inode pointer

- * what does walkpgdir do:
Now the argument given to walkpgdir is '0'
- * P2V() is done bcoz the address in page table is physical address.

* PGROUNDUP()



→ The invalid page is added only for detecting user stack violation

- * ~~exec()~~ exec() also pushes the argc and argv on the user stack. Bcoz they must be made available to main() of new process